

# Parallel Oblivious Sorting in Intel SGX Enclave

Tianyao Gu, Tian Xie

December 14, 2023

## Abstract

Oblivious algorithms have been a popular research topic in cybersecurity due to its resistance to side channel attacks. We implemented a parallel oblivious sorting algorithm in C++ using Intel SGX Enclave, a hardware-based trusted-computing environment. We leveraged OpenMP to achieve multithreading on a 36-core CPU and further improved parallelism through SIMD. We parallelized both the in-enclave computation and the communication with untrusted memory. Our implementation achieves a speedup of 7.5x over the original serial implementation with 8 threads and 16x with 32 threads.

## Background of Intel SGX

Intel SGX is a hardware-based trusted-computing environment, which provides a protected region of memory, known as the Enclave Page Cache (EPC). EPC has limited size, and the communication between EPC and untrusted memory is time-consuming. Therefore, it is important to minimize the amount of data swapped in and out of EPC.

## Background of Oblivious Sorting Algorithms

Oblivious algorithms ensure that memory access and page swap patterns are independent of secret data, thereby resists side channel attacks.

## Background of Flex-way Butterfly Oblivious Sort

To achieve oblivious sorting, Flex-way Butterfly Oblivious Sort applies a random shuffling algorithm followed by a non-oblivious comparison-based sorting algorithm. The key component of the shuffling algorithm is a multi-way butterfly network, emulated with a building block called Merge-split. The algorithm ensures negligible overflow probability and achieves optimal complexity in terms of both computation and I/O.

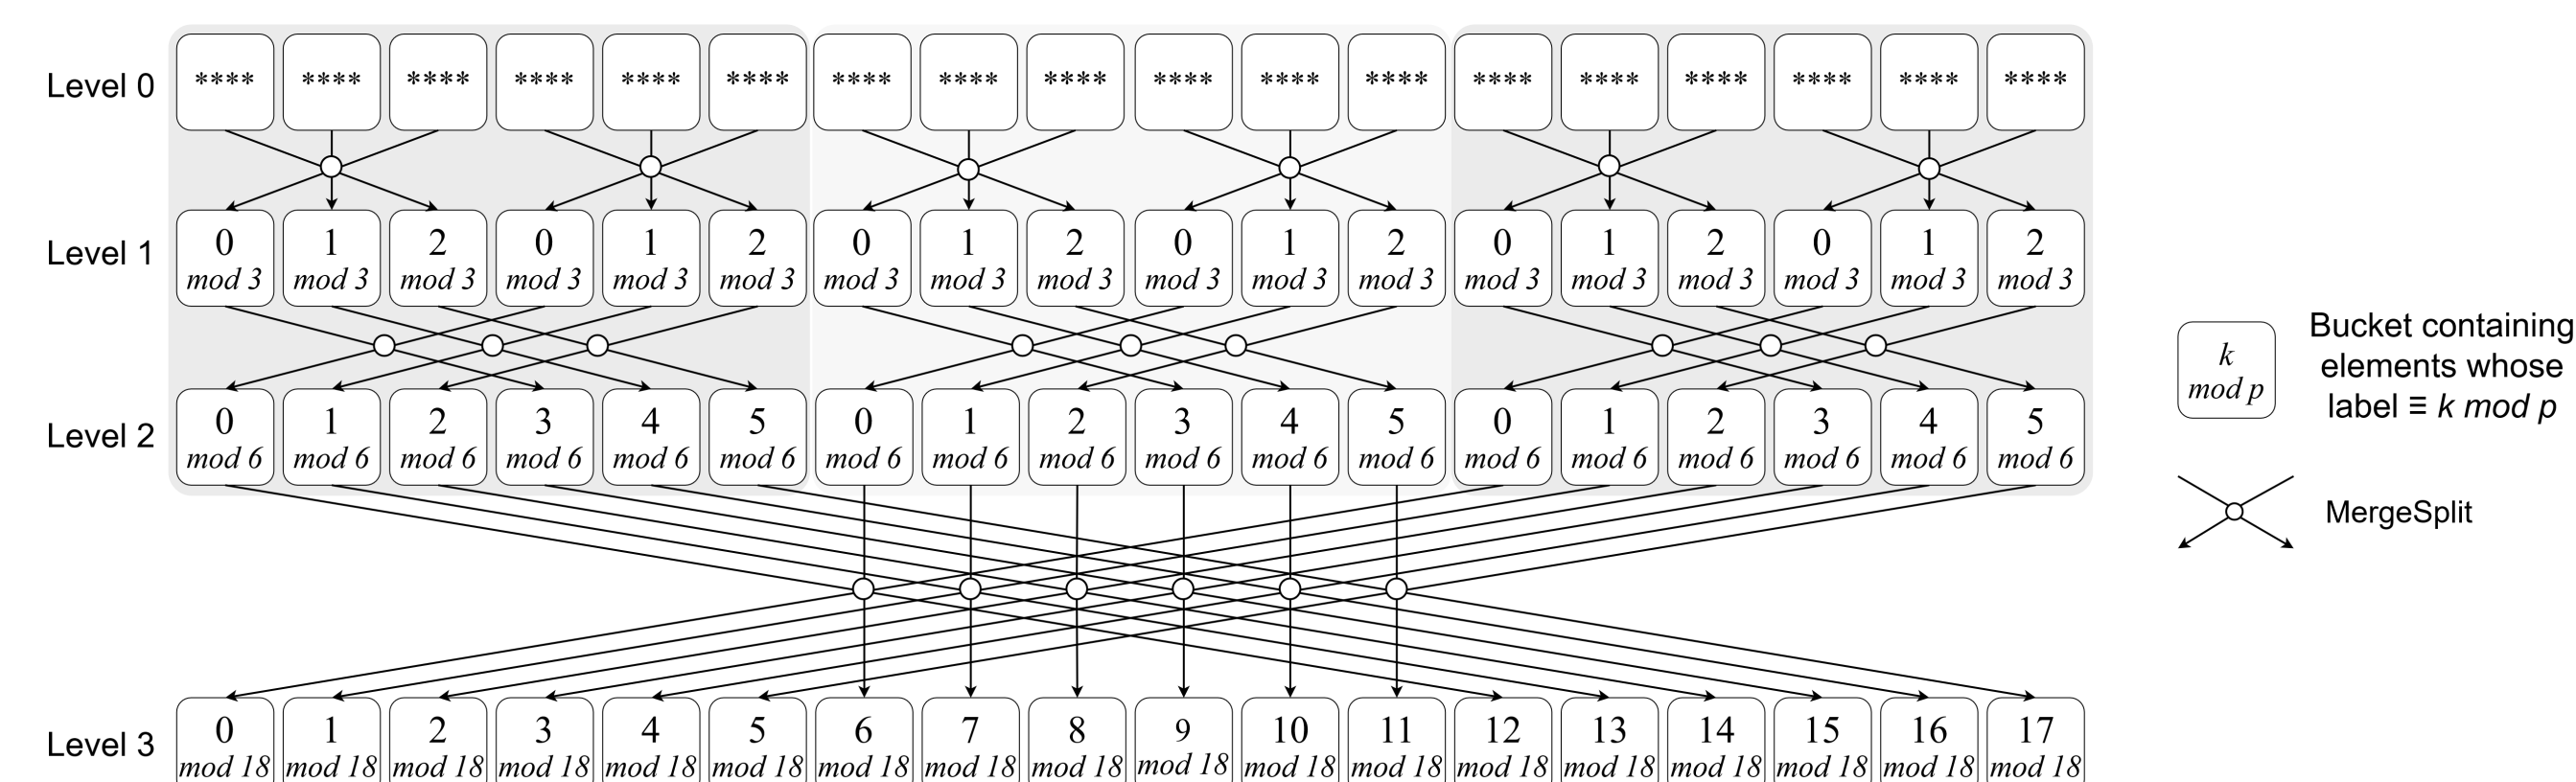


Figure: A Multi-way Butterfly Network

## Parallelism Abstraction

- ▶ Across merge-split task in each batch
- ▶ Within each merge-split task
- ▶ Parallel External Merge Sort
- ▶ Parallel I/O

## Parallelism in I/O

- ▶ **Parallelism in I/O:** Identified three I/O patterns and used different strategies for parallelization.
- ▶ **Parallelism in Element Movement:** Utilized SIMD for parallelizing data movement, especially for elements with large payloads.
- ▶ **Parallelism in Pseudo-Random Numbers Generation:** Parallel generation of secure pseudo-random numbers using multiple generators.

## Implementation

- ▶ **Overview:** Implemented algorithm based on open-source code [5]. Modifications for parallelism and optimizations for high-end servers and consumer-grade processors.
- ▶ **Parallelize Butterfly Network:** Applied OpenMP directives for parallelization, changed routing schedule to iterative, and optimized parameters for maximum parallelism.
- ▶ **Parallelize I/O:** Identified three I/O patterns and used different strategies for parallelization.
- ▶ **Parallelize Element Movement:** Utilized SIMD for parallelizing element movement, applying C++ intrinsics with AVX512, AVX2, and SSE2 instruction sets.
- ▶ **Parallelize Pseudo-Random Number Generation:** Implemented multiple pseudo-random number generators for parallel generation.
- ▶ **Optimize Memory Utilization:** Addressed memory fragmentation concerns by allocating a memory pool and implementing a custom allocator.

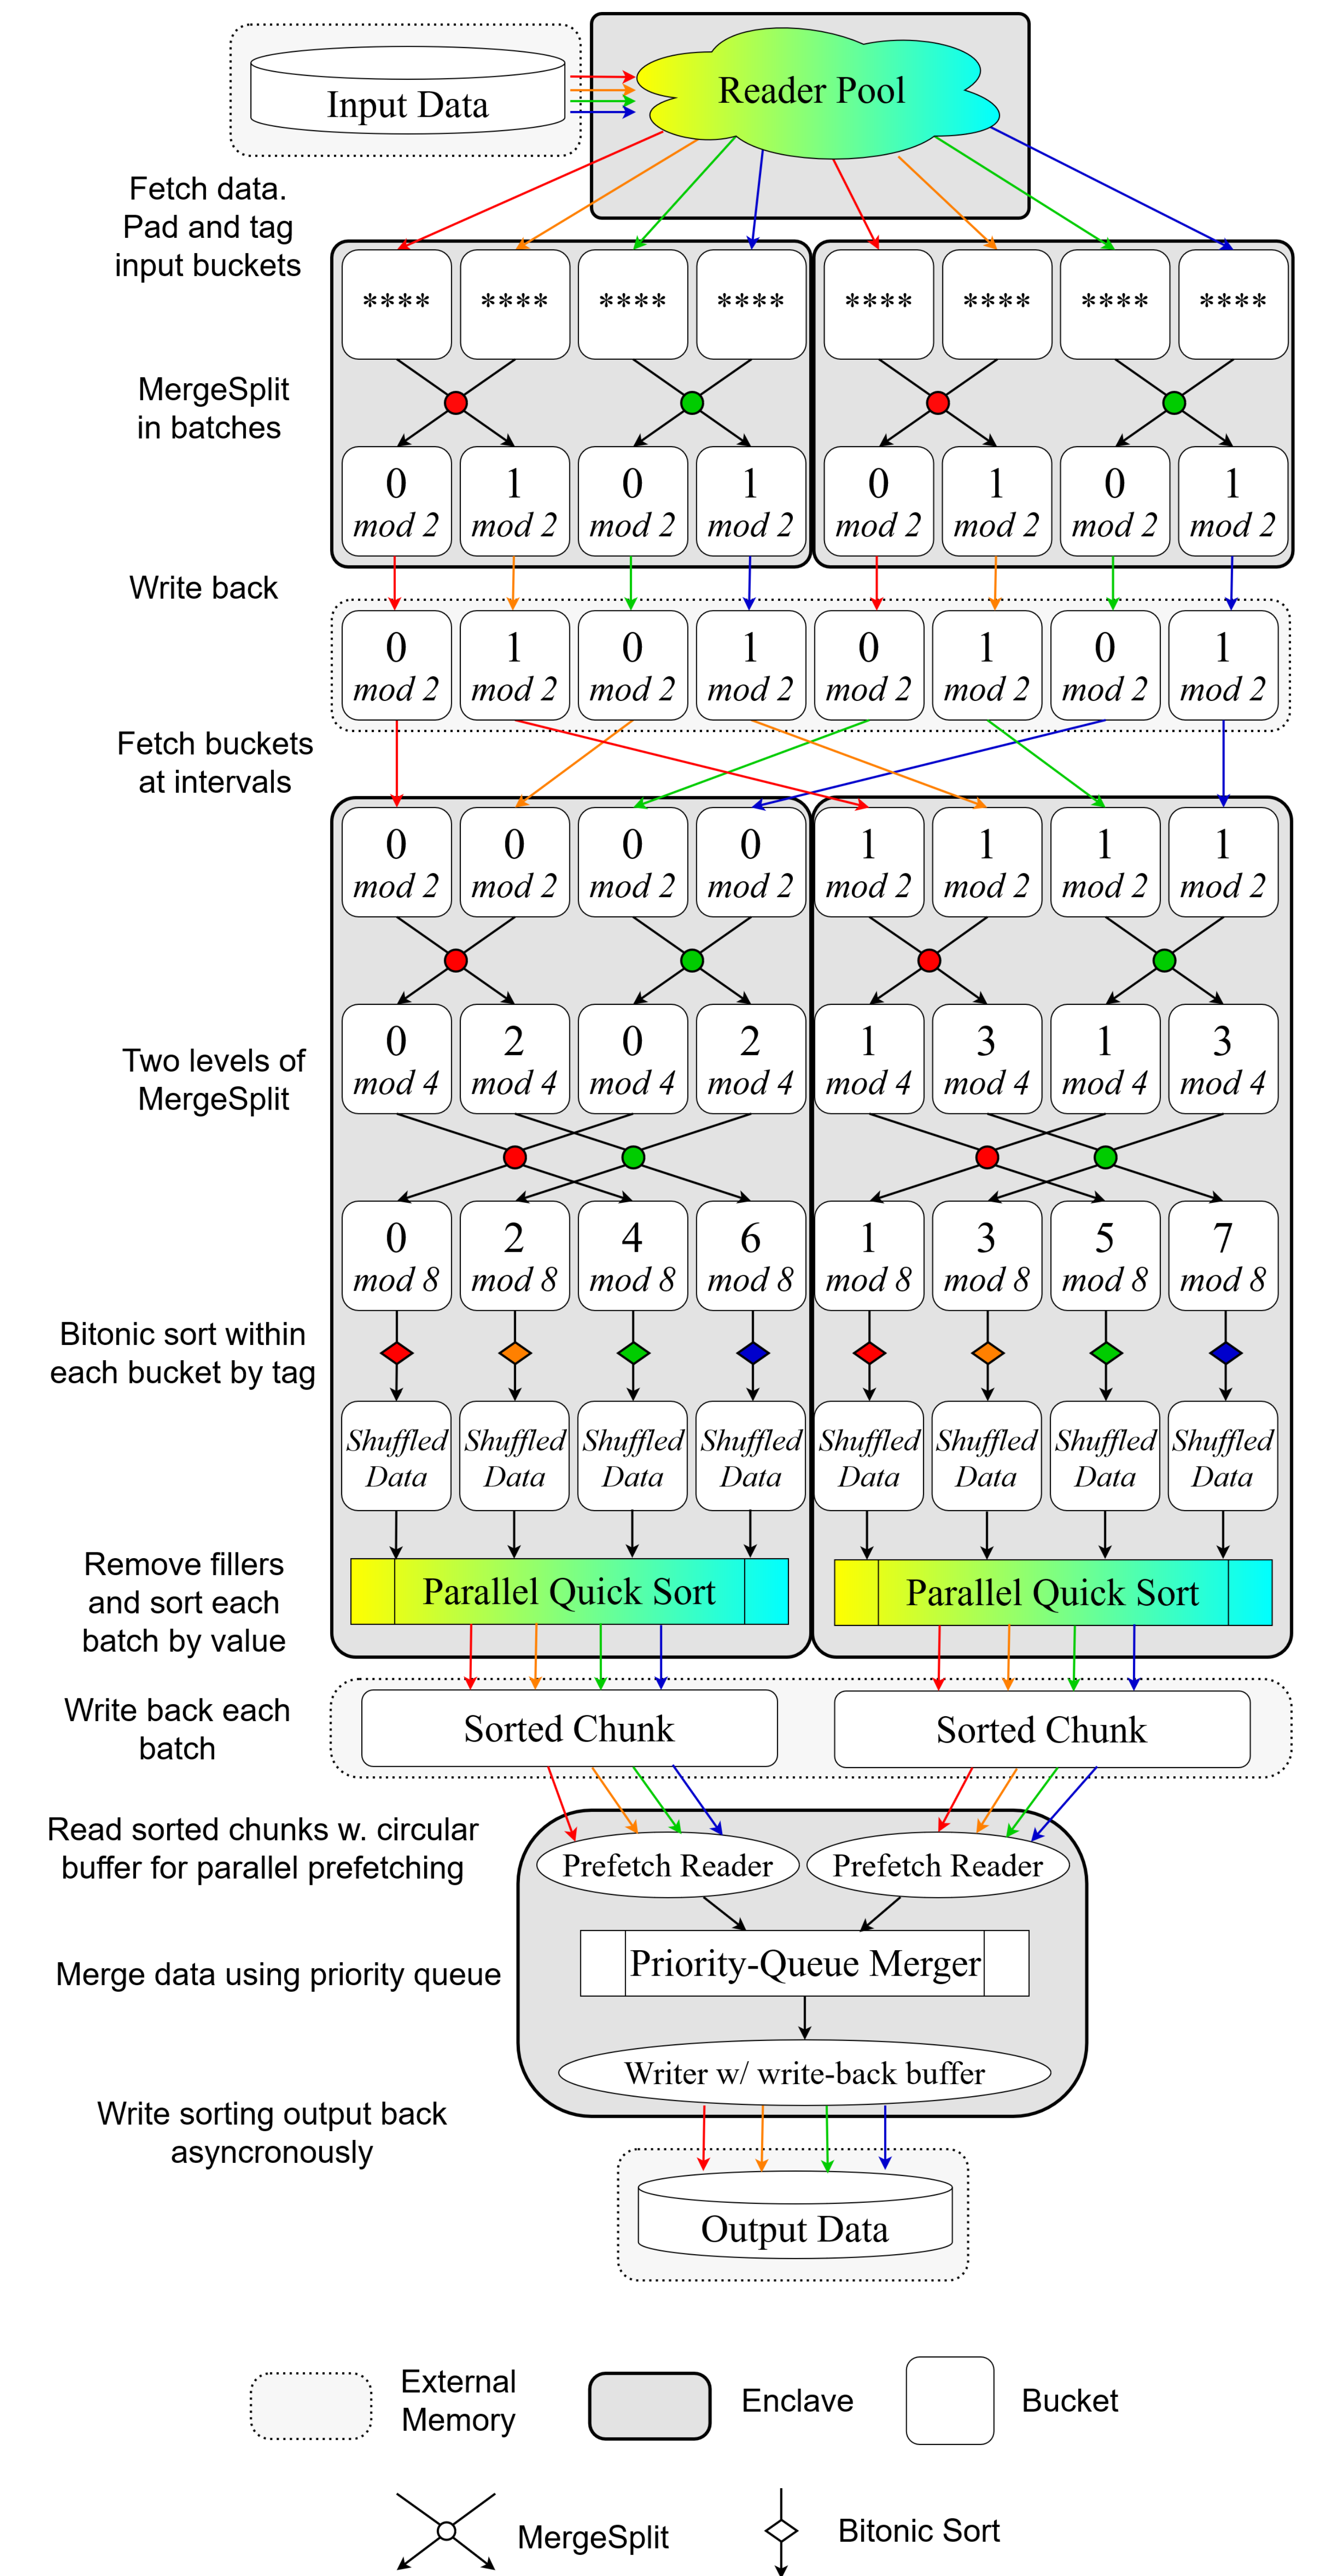


Figure: Diagram of Parallel Oblivious Sort on a toy example. Each color represents a thread.